



# Filtering based on process ranking

Mahsa Shakeri

May 06, 2019

Polytechnique Montréal  
DORSAL Lab

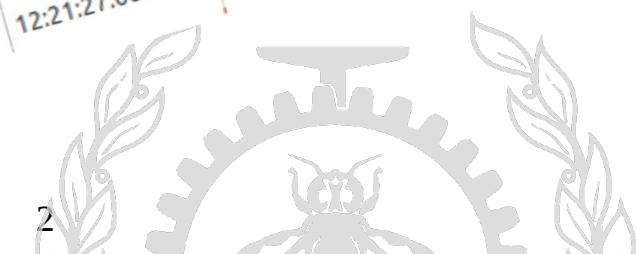
# Introduction

---

Tracing is used for performance analysis and debugging.

Usually the trace size is big and we need some heuristics to cope with it.

Process	TID	PTID	Birth time
gnome-pty-helper	3609	3596	18:15:38.7695138
▼ bash	3610	3596	18:15:38.7695147
ssh	8972	3610	18:1!
▼ bash	5332	3596	18:1!
ls	15298	5332	18:1!
▼ bash	10234	3596	18:1!
ssh	8989	10234	18:1!
▼ bash	11275	3596	18:1!
ssh	15266	11275	18:1!
▼ bash	15291	15266	18:1!
▼ bash	15292	15291	18:1!
sudo	15294	15266	18:1!
lttng	15302	15266	18:1!
ab	15303	15302	18:1!
▼ sudo	3599	2779	18:1!
lttng	3600	2779	18:1!
dconf worker	3601	2779	18:1!
gmain			
gdbus			
unity-greeter			
llvmpipe-0			
llvmpipe-1			
gmain			
gdbus			
dconf worker			
threaded-ml			
gmain			
gdbus			
lightdm			
gmain			
gdbus			
lightdm			
gmain			
gdbus			
mysqld			
▼ mysqld			
mysqld	1545	861	12:21:27.6052379
mysqld	941	1	12:21:27.6052392
mysqld	955	1	12:21:27.6052408
mysqld	956	1	12:21:27.6052408

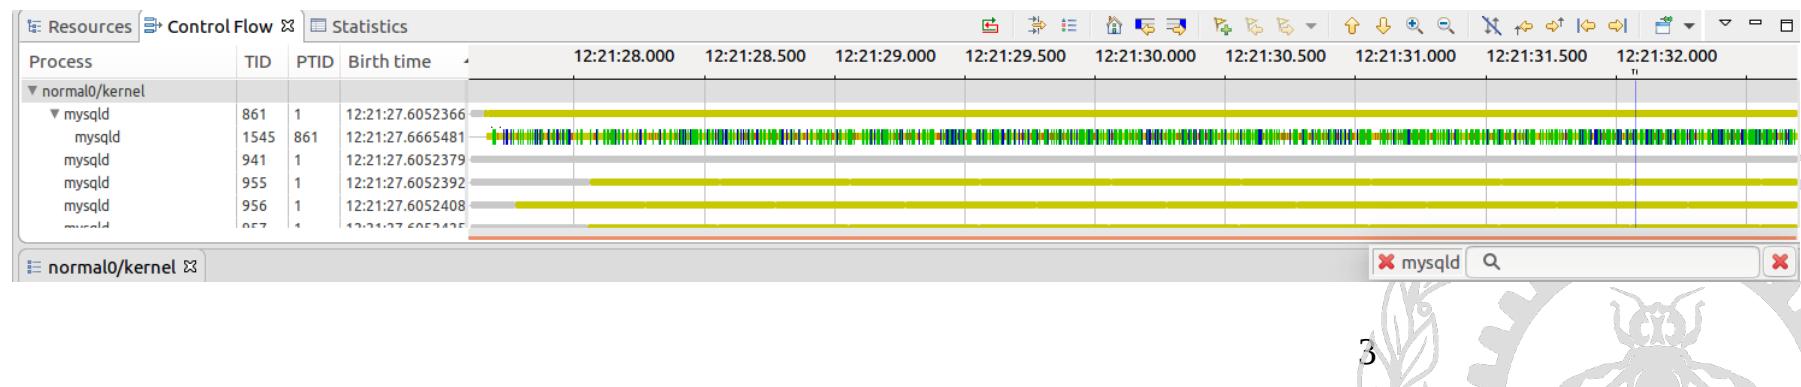


# Current methods



- ✓ Sampling or event filtering.
- ✓ Time view filtering.
- ✓ Global filtering

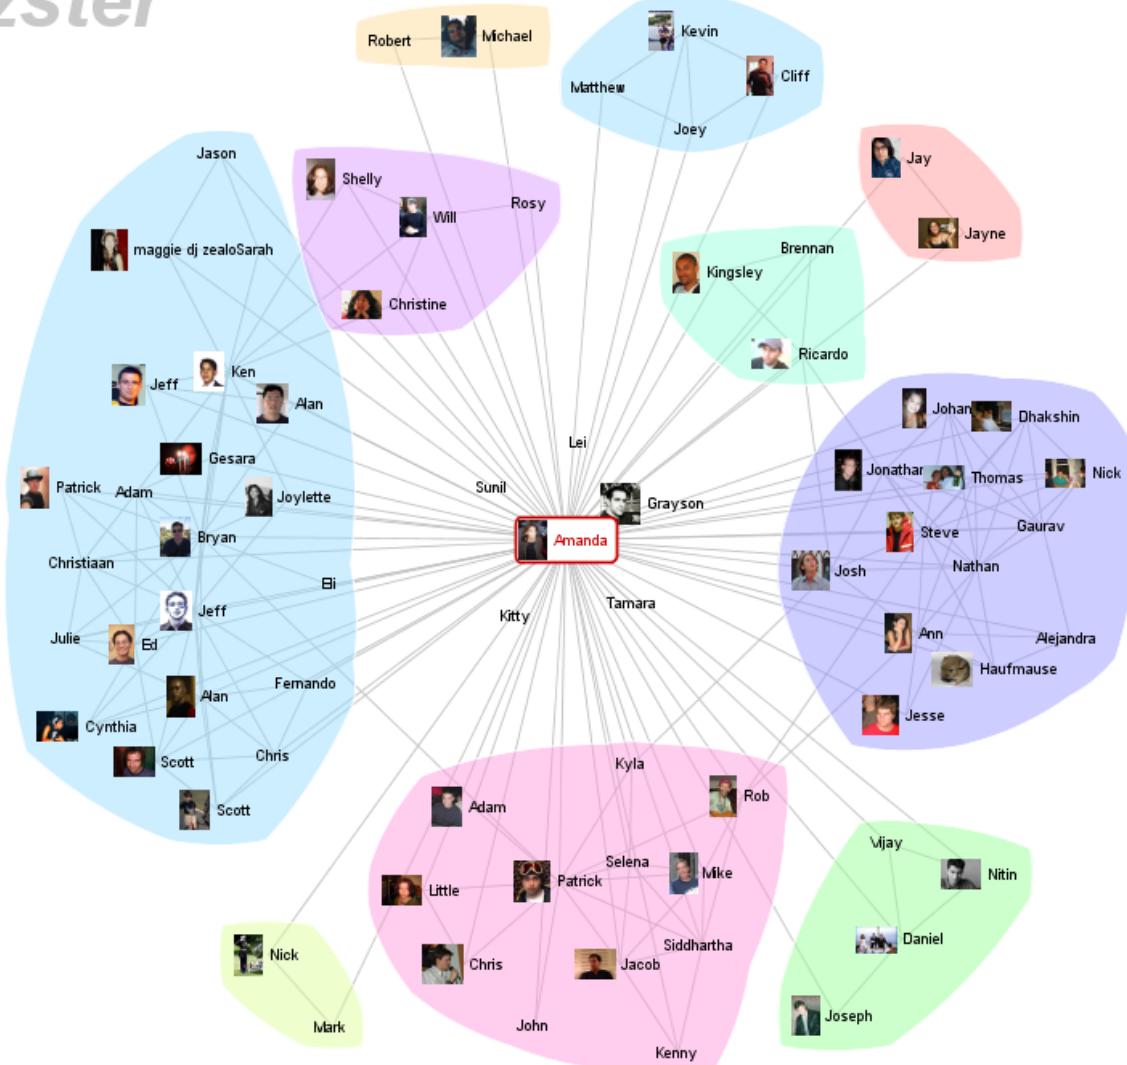
Timestamp	Channel	CPU	Event type	Contents
<srch>	<srch>	<srch>	<srch>	<b>tid=861</b> fd=3, uservaddr=0x7ffe5564c030, addrlen=0x6e, family=1, dport=0, _v4addr_length=0, v4addr=[], _v6addr_length=0, v6addr=[], context...
12:21:27.650 850 349	more-subbuf_0 0		syscall_entry_connect	
12:21:27.650 863 113	more-subbuf_0 0		sched_waking	comm=mysql, tid=861, prio=20, target_cpu=0, context._tid=1544, context._pid=1544, context._procname=sysbench
12:21:27.650 868 663	more-subbuf_0 0		sched_wakeup	comm=mysql, tid=861, prio=20, target_cpu=0, context._tid=1544, context._pid=1544, context._procname=sysbench
12:21:27.650 870 198	more-subbuf_0 0		syscall_exit_connect	ret=0, context._tid=1544, context._pid=1544, context._procname=sysbench
12:21:27.650 873 554	more-subbuf_0 0		sched_switch	prev_comm=sysbench, prev_tid=1544, prev_prio=20, prev_state=0, next_comm=mysql, next_tid=861, next_prio=20, context._tid=1544, context._pid=1544, context._procname=sysbench
12:21:27.650 886 455	more-subbuf_0 0		syscall_exit_poll	ret=1, nfds=2, fds_length=1, fds=[[fd=23, raw_events=0x1, events_POLLIN=1, events_POLLPRI=0, events_POLLOUT=0, events_POLLFD=0], [fd=23, upeer_sockaddr=0x7fff9470d50, upeer_addrlen=128, context._tid=861, context._pid=861, context._procname=mysql]
12:21:27.650 895 764	more-subbuf_0 0		syscall_entry_accept	fd=23, upeer_sockaddr=0x7fff9470d50, upeer_addrlen=128, context._tid=861, context._pid=861, context._procname=mysql
12:21:27.650 904 076	more-subbuf_0 0		syscall_exit_accept	ret=39, upeer_addrlen=2, family=1, sport=0, _v4addr_length=0, v4addr=[], _v6addr_length=0, v6addr=[], context._tid=861, context._pid=861, context._procname=mysql
12:21:27.650 918 648	more-subbuf_0 0		syscall_entry_futex	uaddr=31402148, op=133, val=1, utime=1, uaddr2=31402144, val3=67108865, context._tid=861, context._pid=861, context._procname=mysql
12:21:27.650 923 327	more-subbuf_0 0		sched_waking	comm=mysql, tid=1291, prio=20, target_cpu=0, context._tid=861, context._pid=861, context._procname=mysql
12:21:27.650 926 771	more-subbuf_0 0		sched_wakeup	comm=mysql, tid=1291, prio=20, target_cpu=0, context._tid=861, context._pid=861, context._procname=mysql
12:21:27.650 928 588	more-subbuf_0 0		syscall_exit_futex	ret=1, uaddr=31402148, uaddr2=31402144, context._tid=861, context._pid=861, context._procname=mysql



# Idea

## Graph theory

vizster



Picture from: <http://vis.stanford.edu/jheer/projects/vizster/>

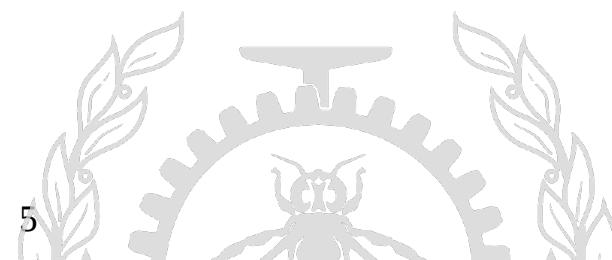
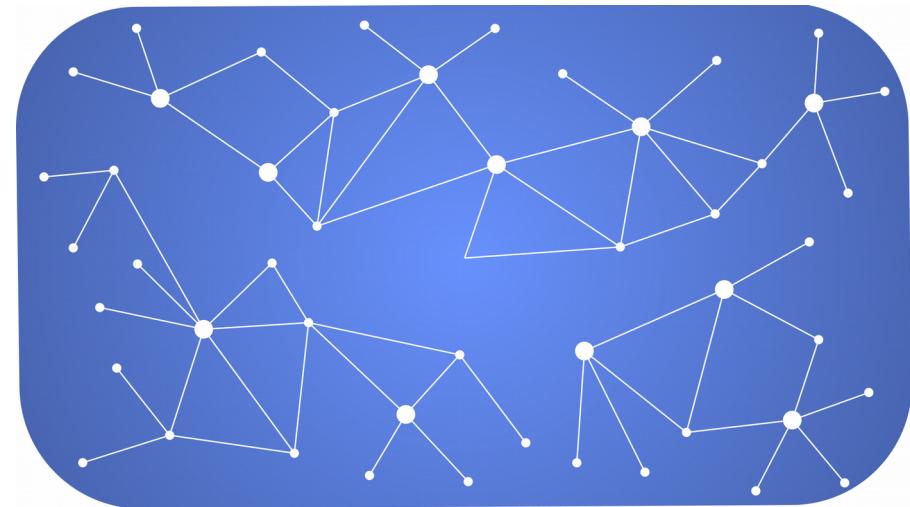


# Method

---

## I.

- ✓ Extract the processes interactions based on the execution graph.
- ✓ Building a graph  $G = (V, E)$
- ✓  $V$  includes all the processes running on the system
- ✓  $E$  is a set of edges defined based on the interactions between the processes.

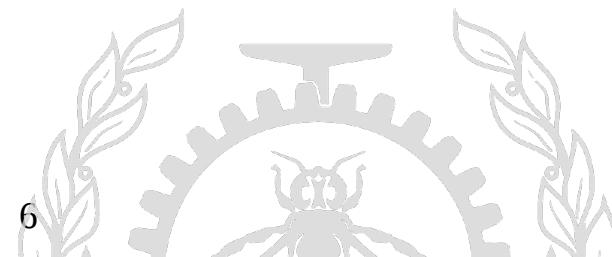
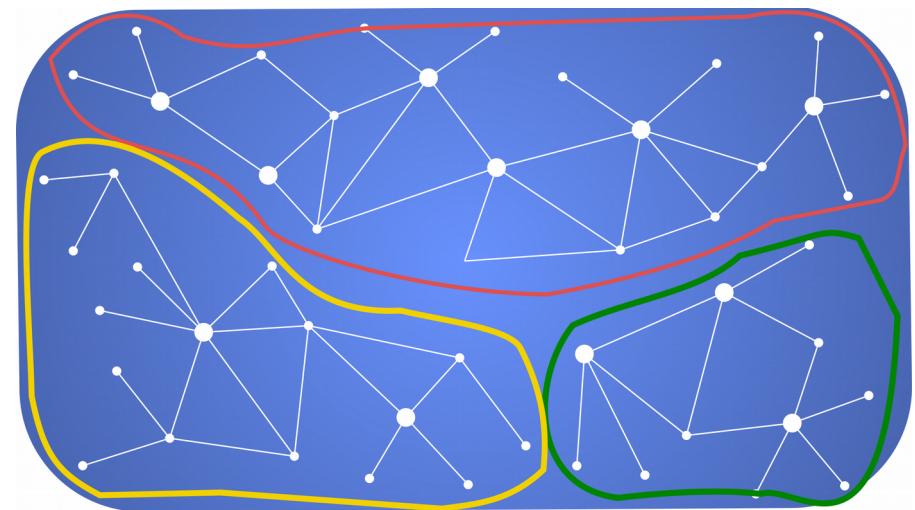


# Method

---

## II.

Compute connected components on the graph to extract all the subgraphs. (e.g., BFS)

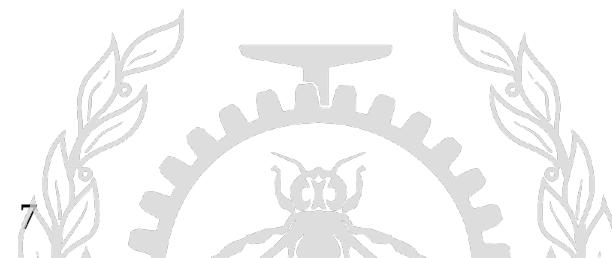
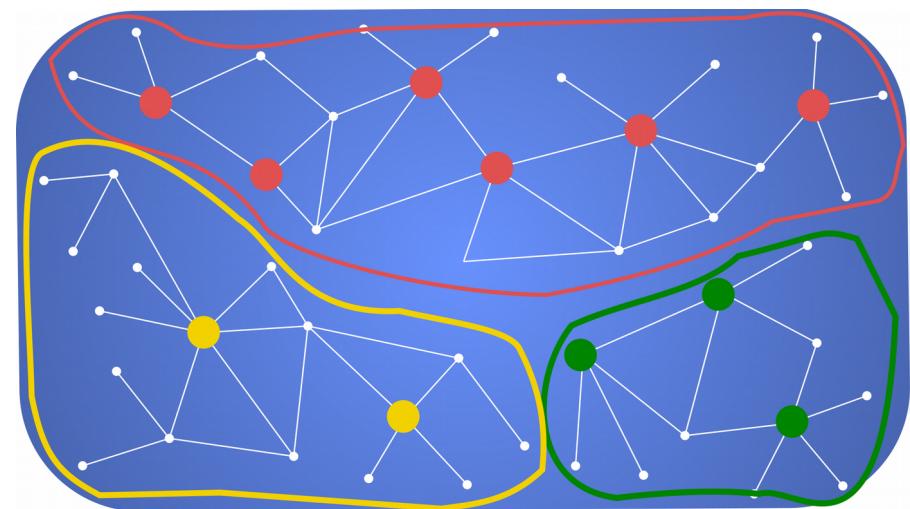


# Method

---

## III.

- ✓ Apply a ranking algorithm on each connected component to assign a rank to each process
- ✓ Choose the top-K processes..



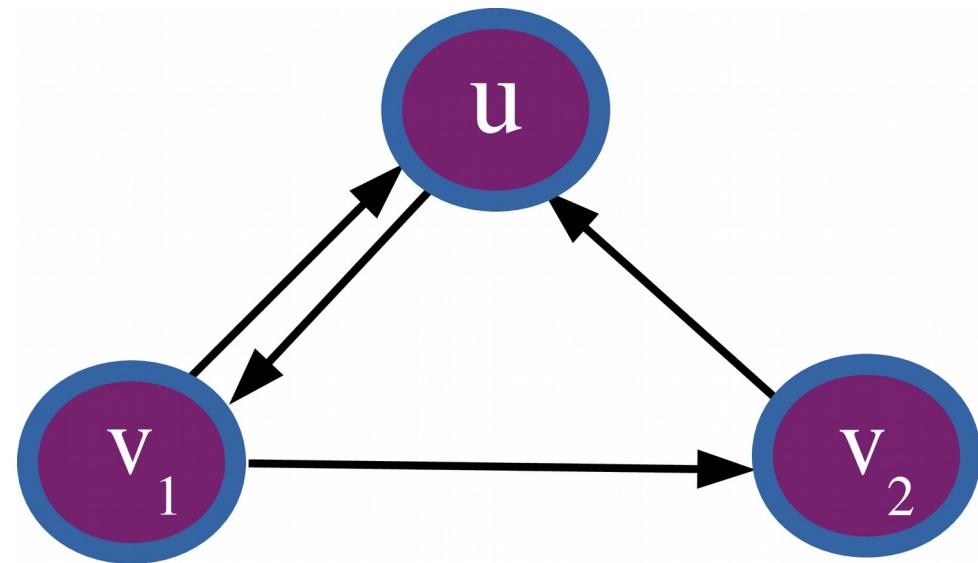
# Method

---

**Modified PageRank** algorithm: Includes only **one** iteration

Node rank initialization: equal weights, or even CPU/IO usage

$$R(u) = c \sum_{v \in B_u} \frac{R(v)}{N_v}$$

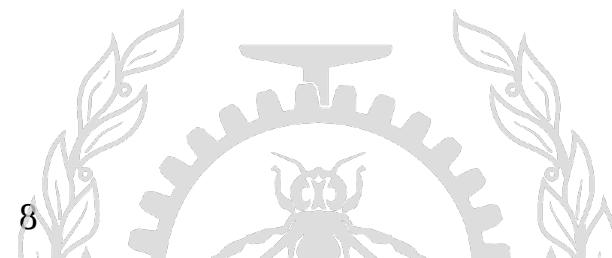


u: a process

$B_u$ : the set of u's in-edges

$N_v$ : the number of out-edges of process v

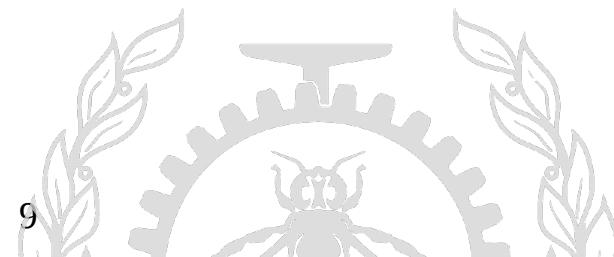
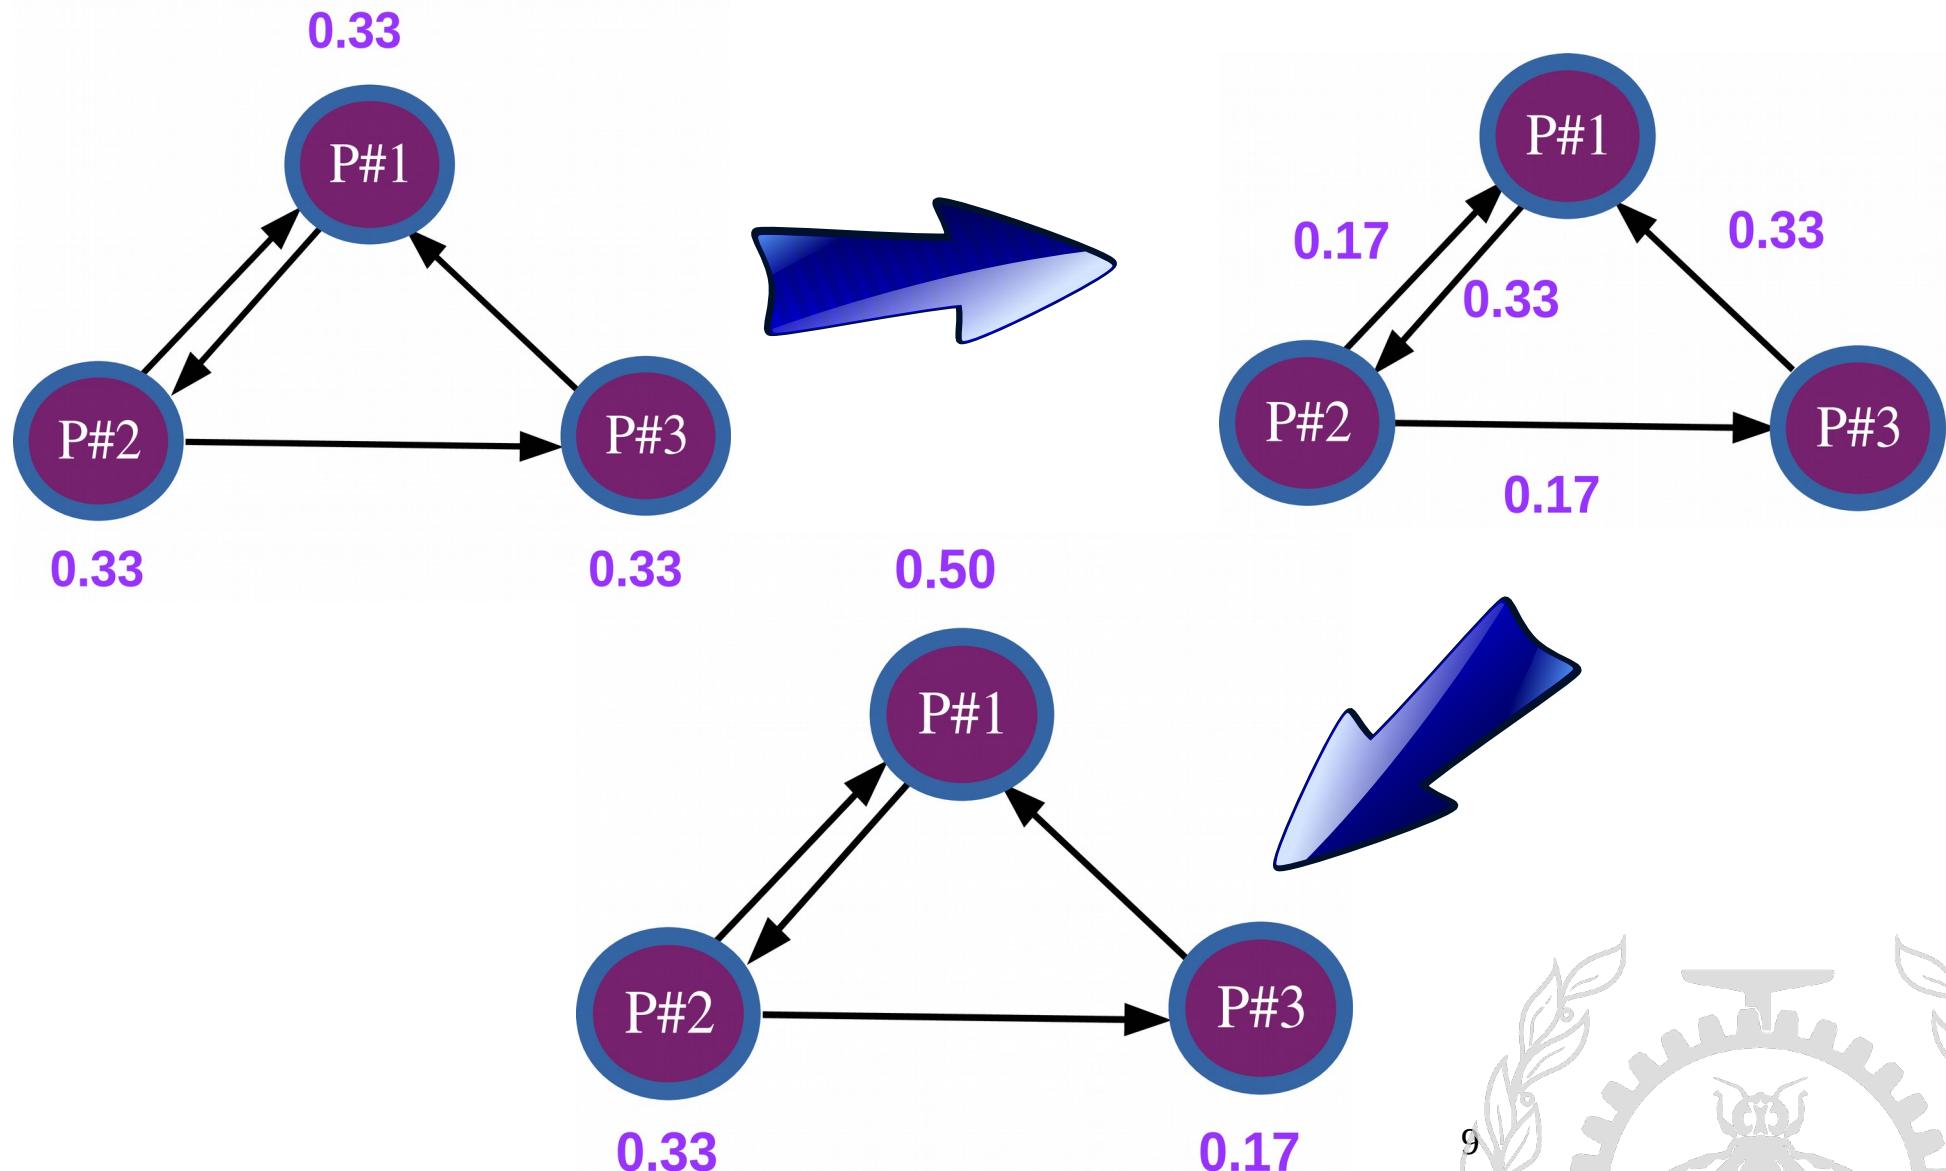
c: the normalization factor



# Method

---

Modified PageRank algorithm:

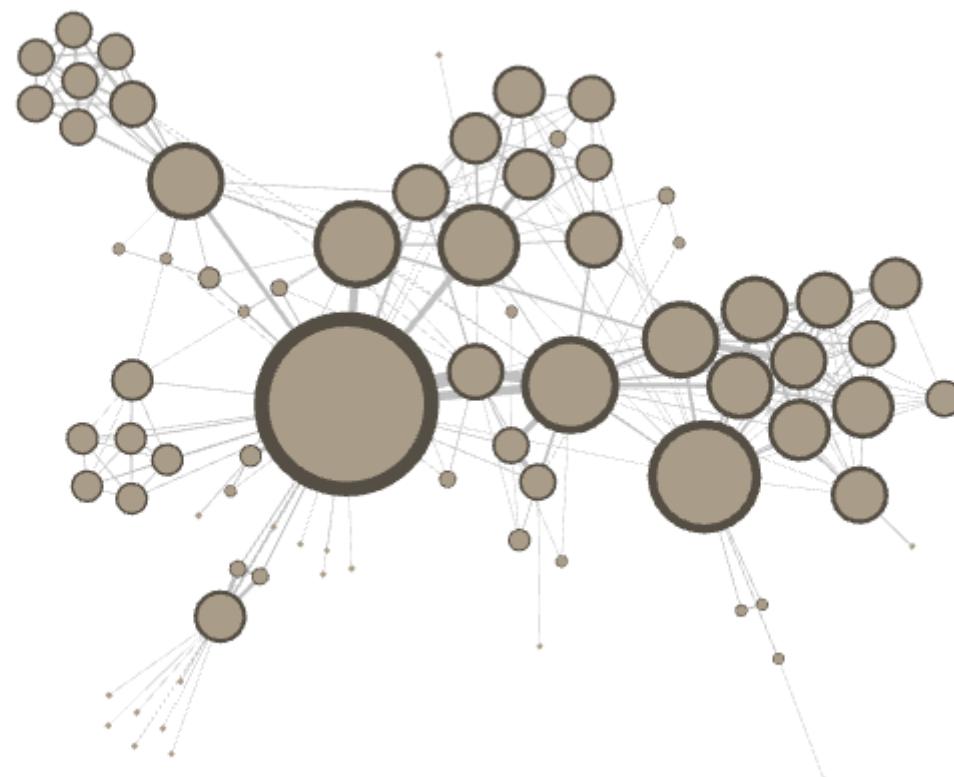


# Method

---

NB. Computing connected components on the graph

Other approaches:  
e.g., Graph Modularity

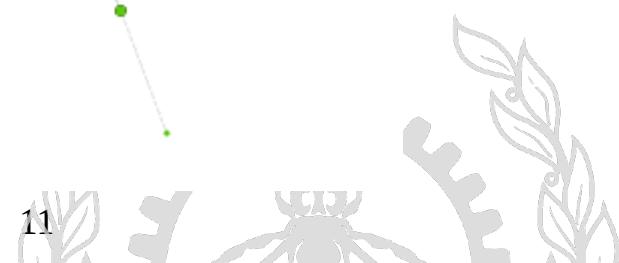
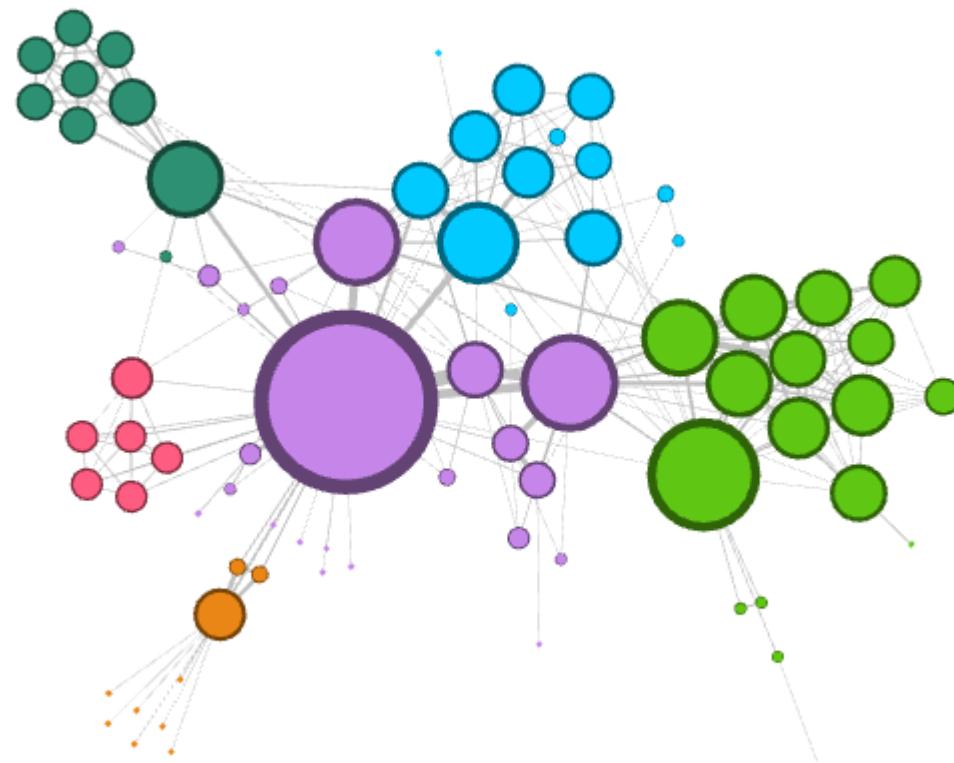


# Method

---

NB. Computing connected components on the graph

Other approaches:  
e.g., Graph Modularity



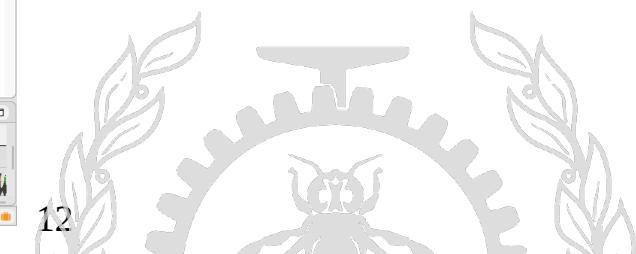
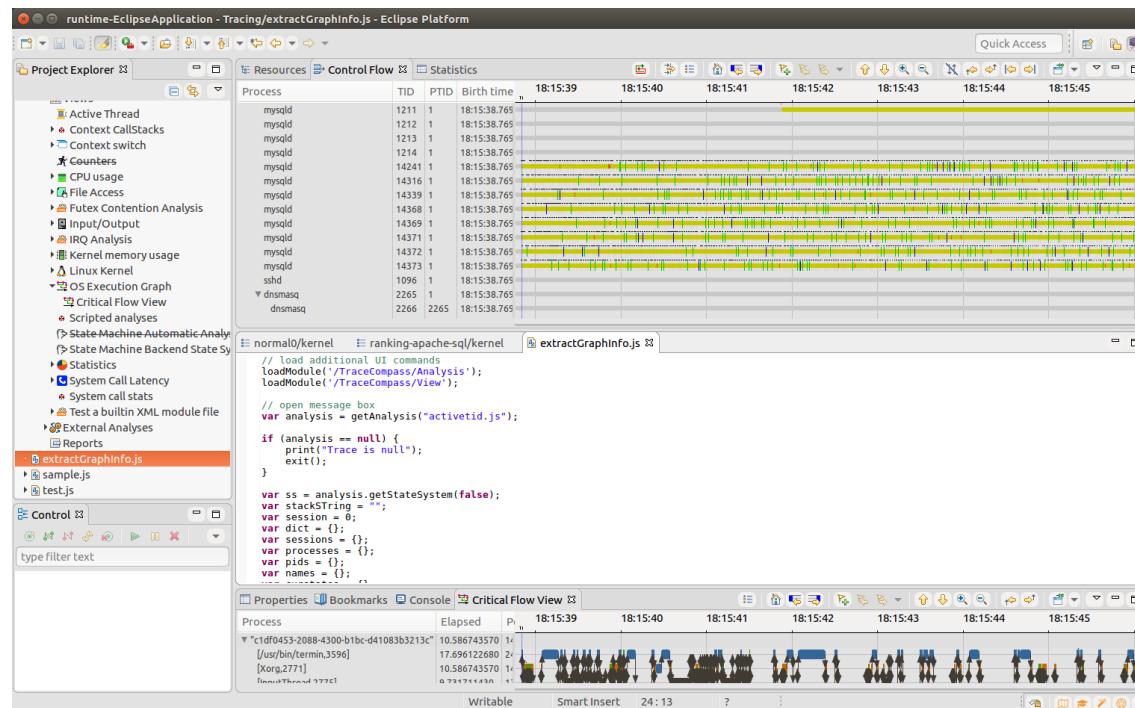
# Use case I

Trace Compass:

Extract process interactions using EASE scripting.

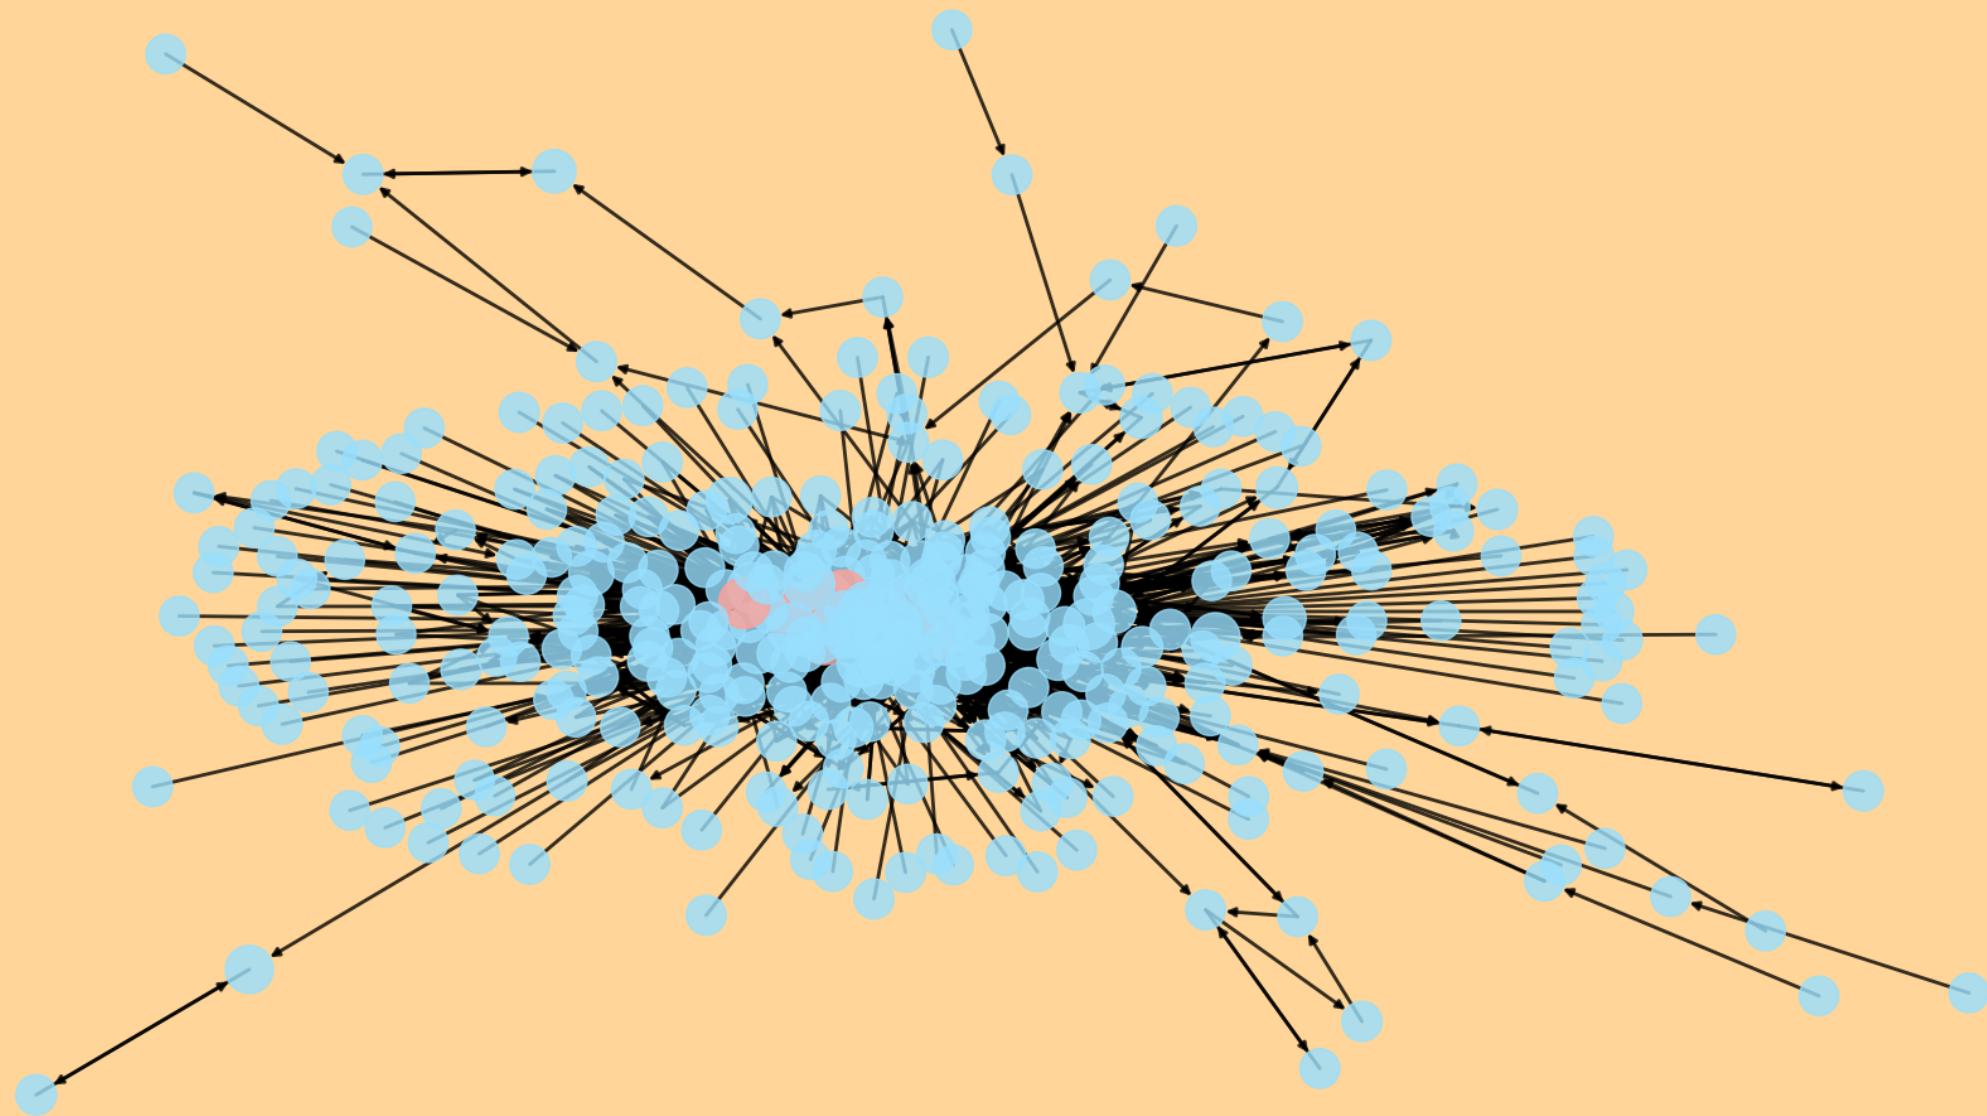
Providing a view showing top-ranked processes in each subgraph.

Filtering top-ranked processes in event view.



# Use case I

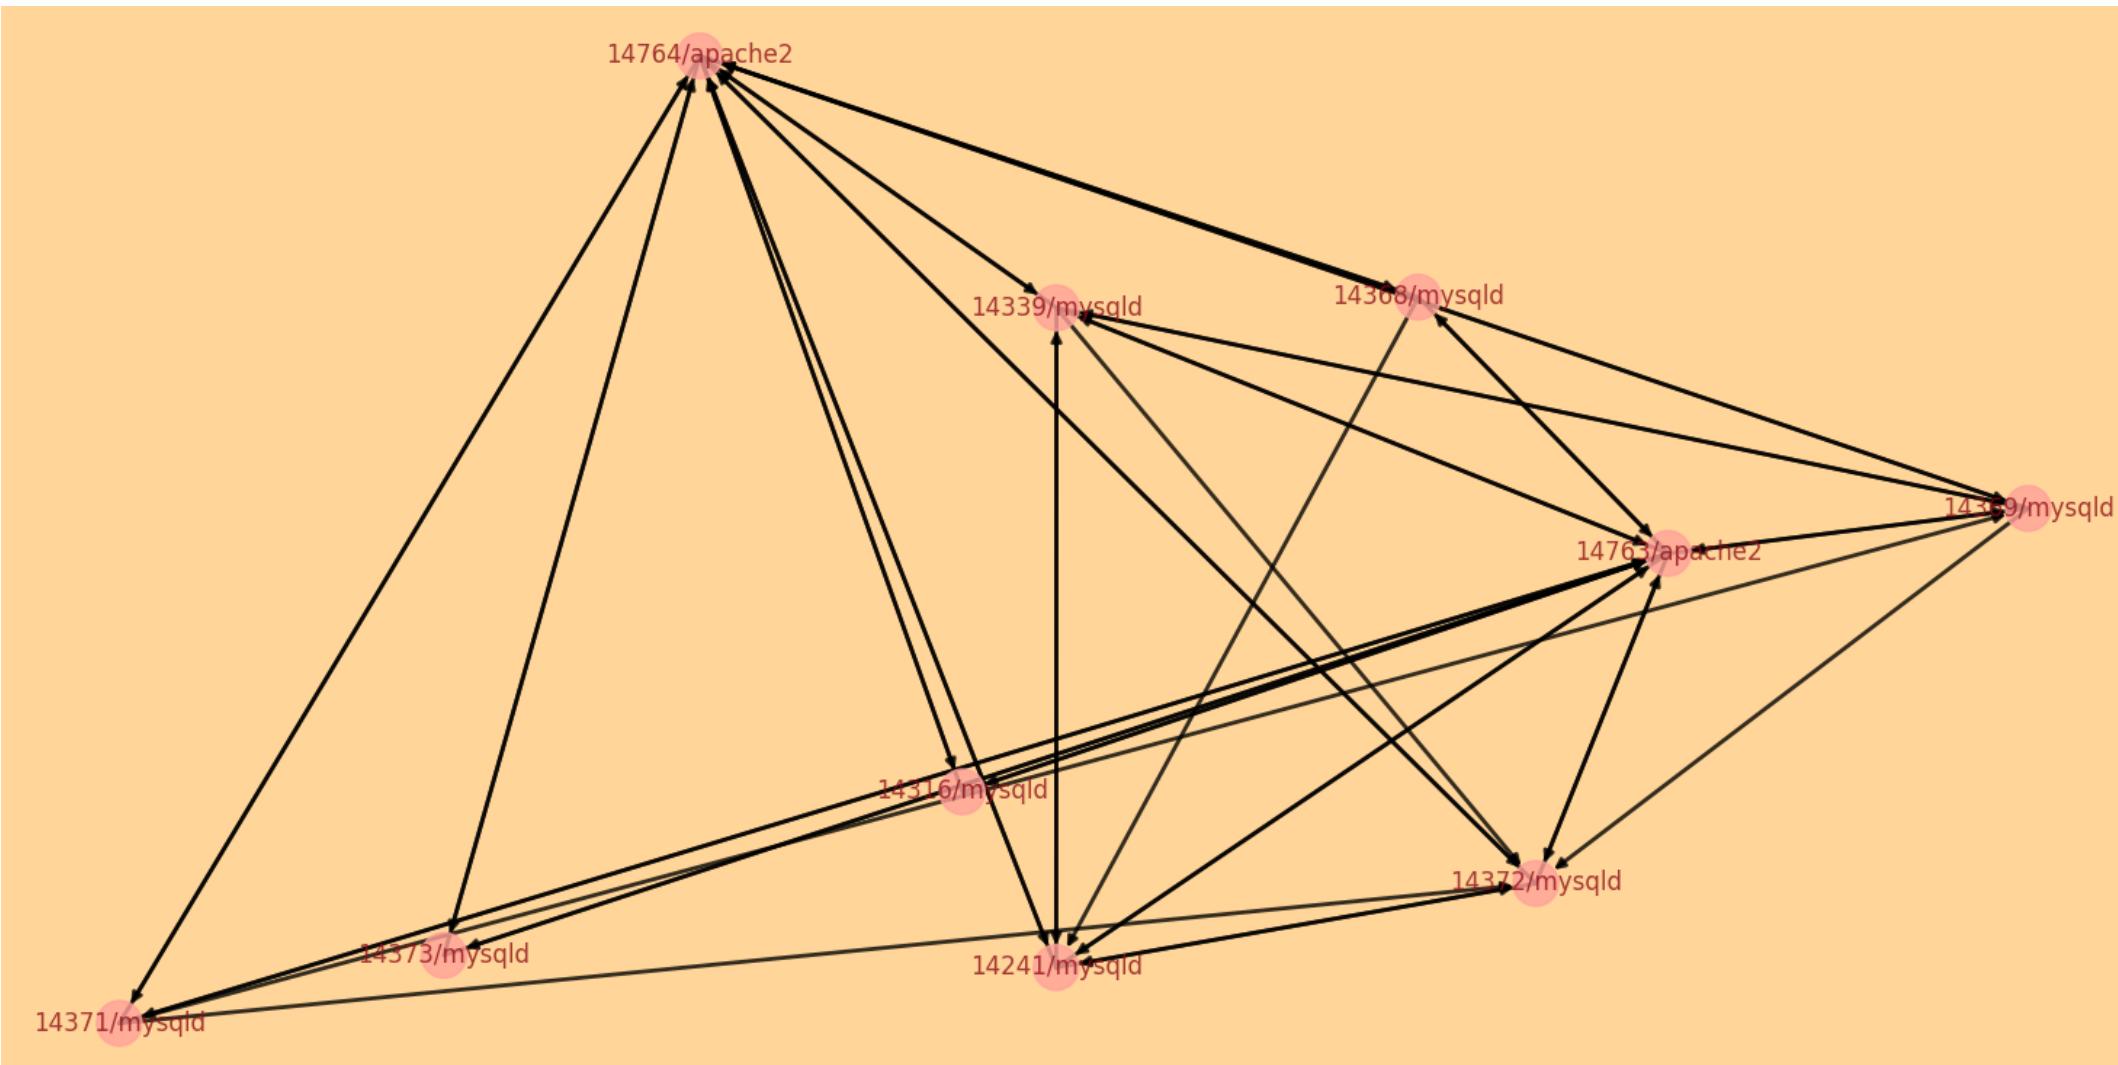
---



# Use case I

---

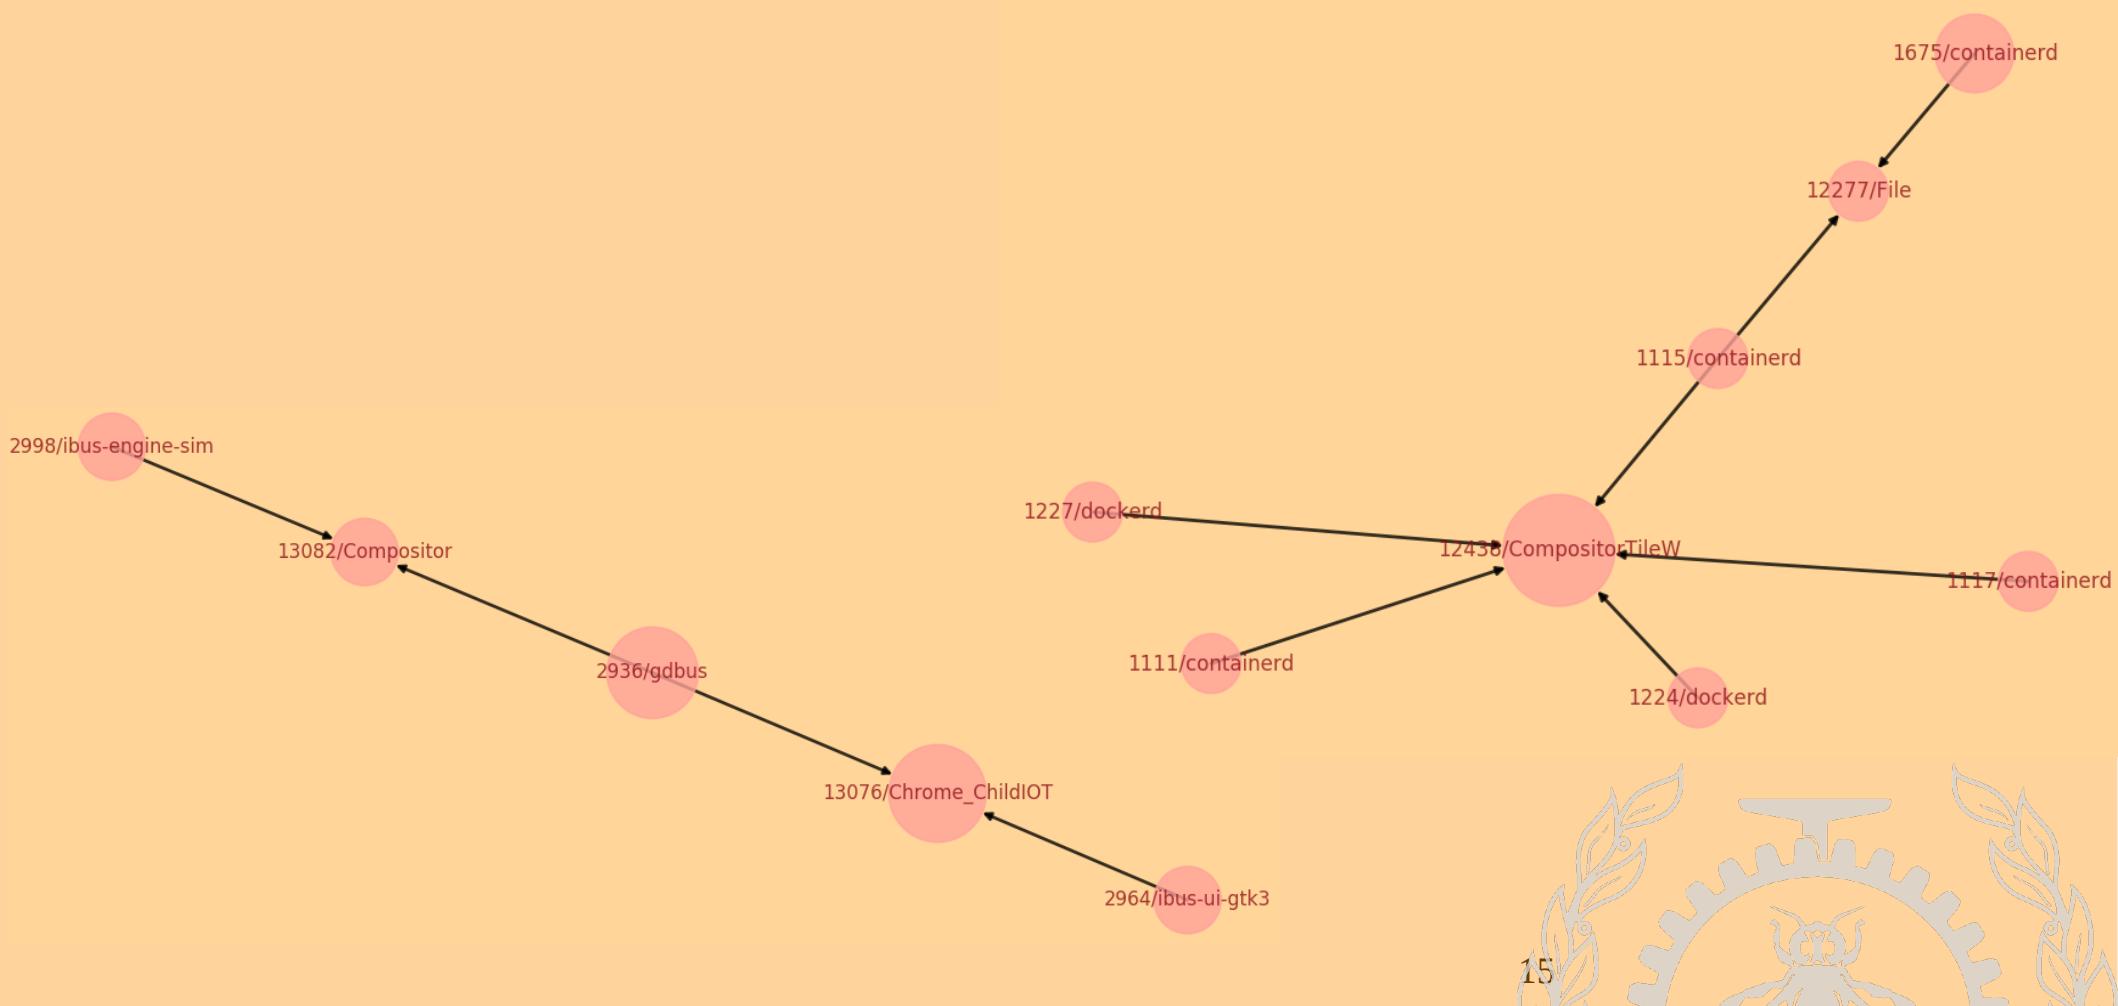
Top-ranked processes in subgraph I



# Use case I

---

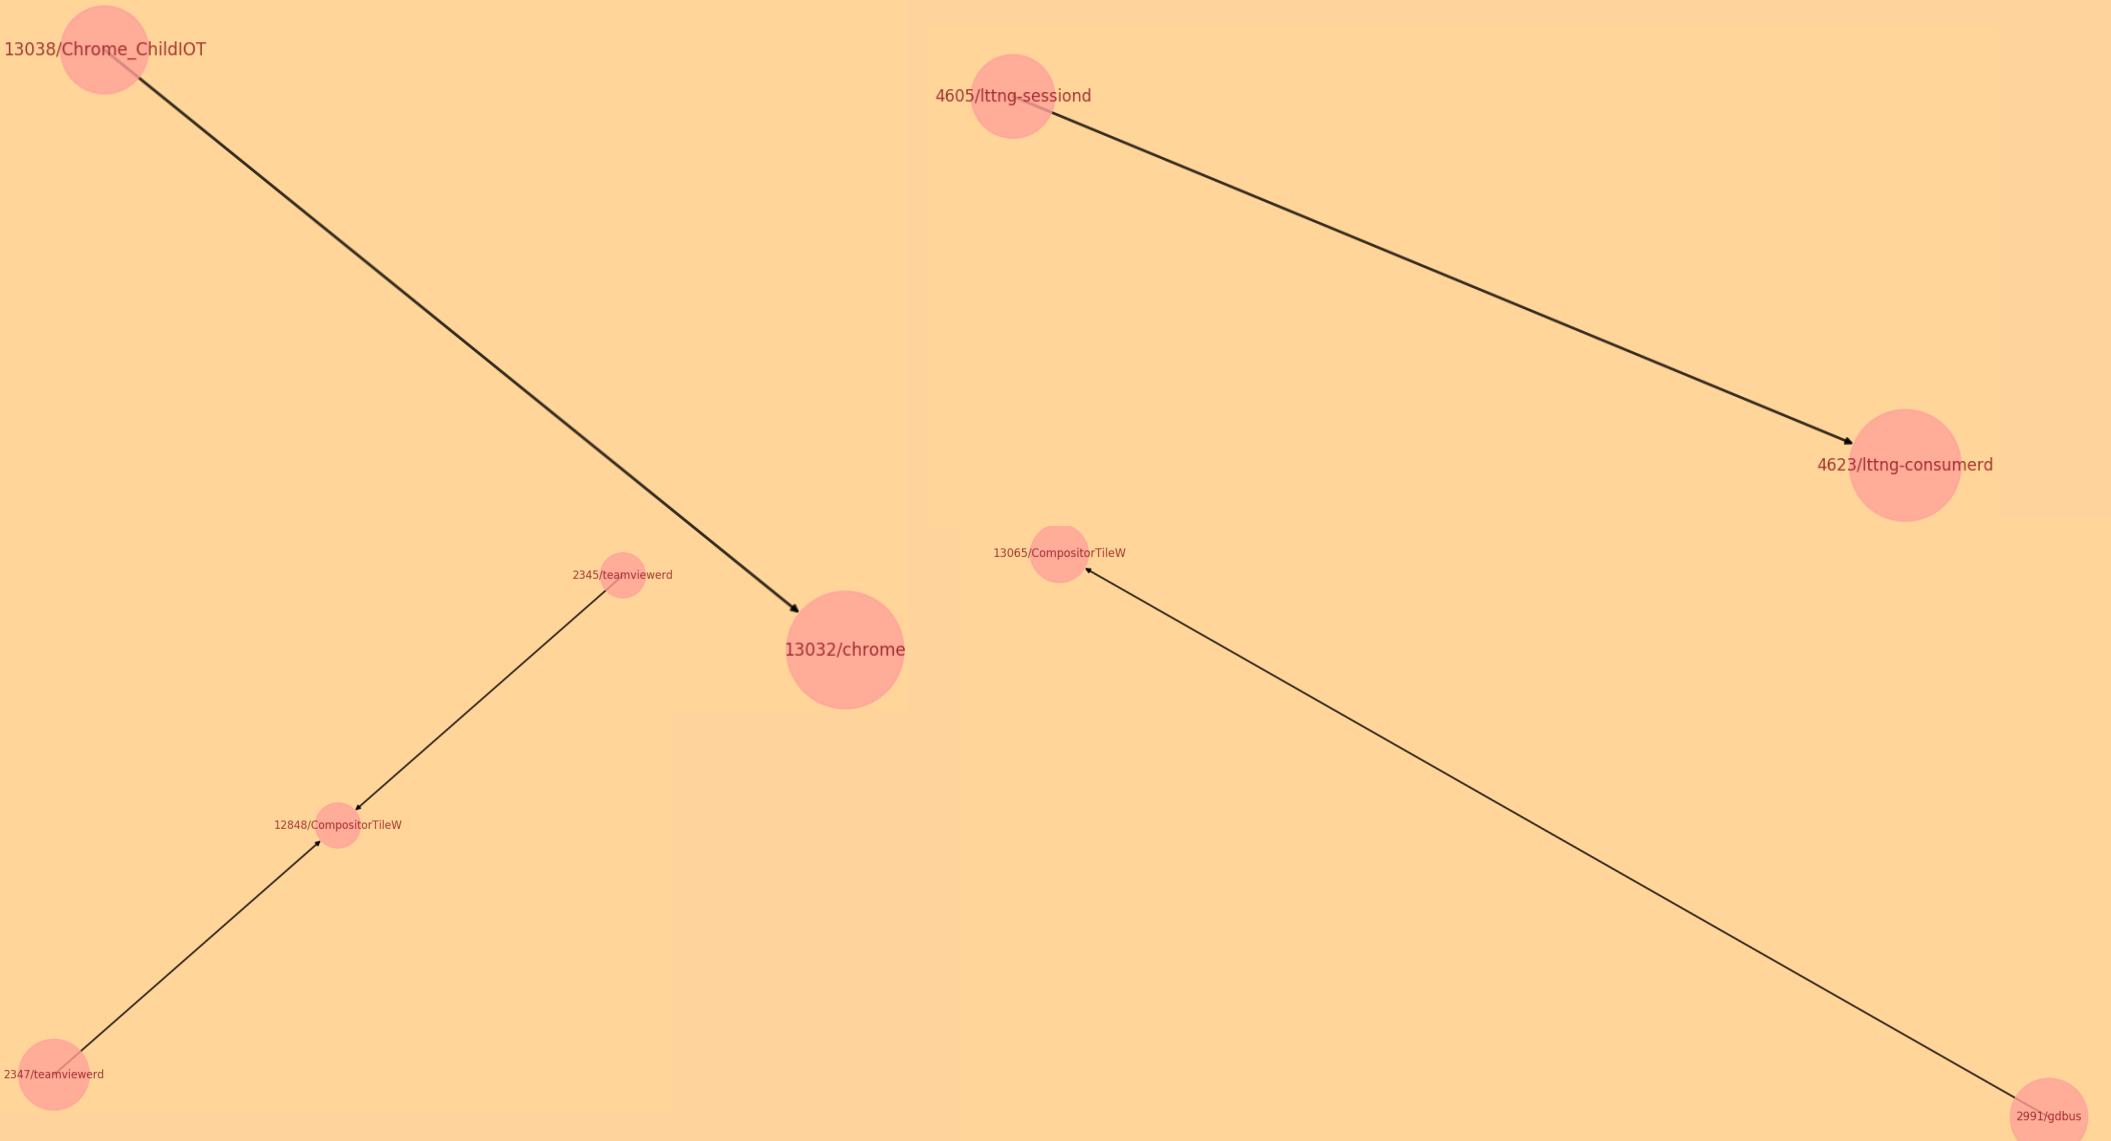
Top-ranked processes in subgraph II and III



# Use case I

---

Top-ranked processes in other subgraphs:

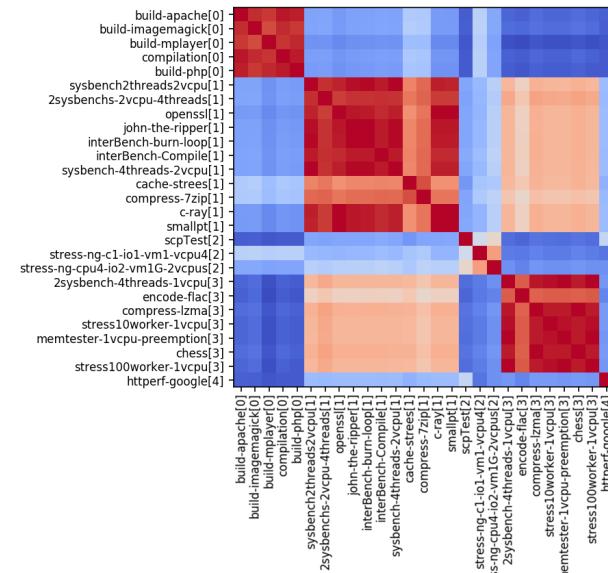
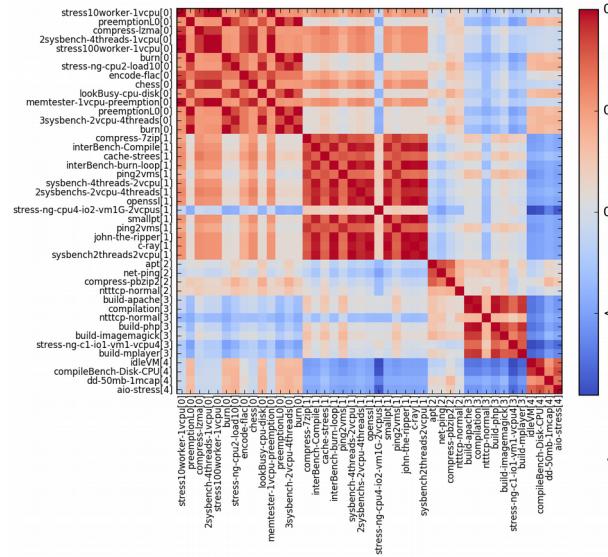


# Use case II

## VM clustering<sup>1</sup>:

Identifying the top important processes running on each VM.

Filtering out the irrelevant processes to provide a more distinct VM workload clustering.

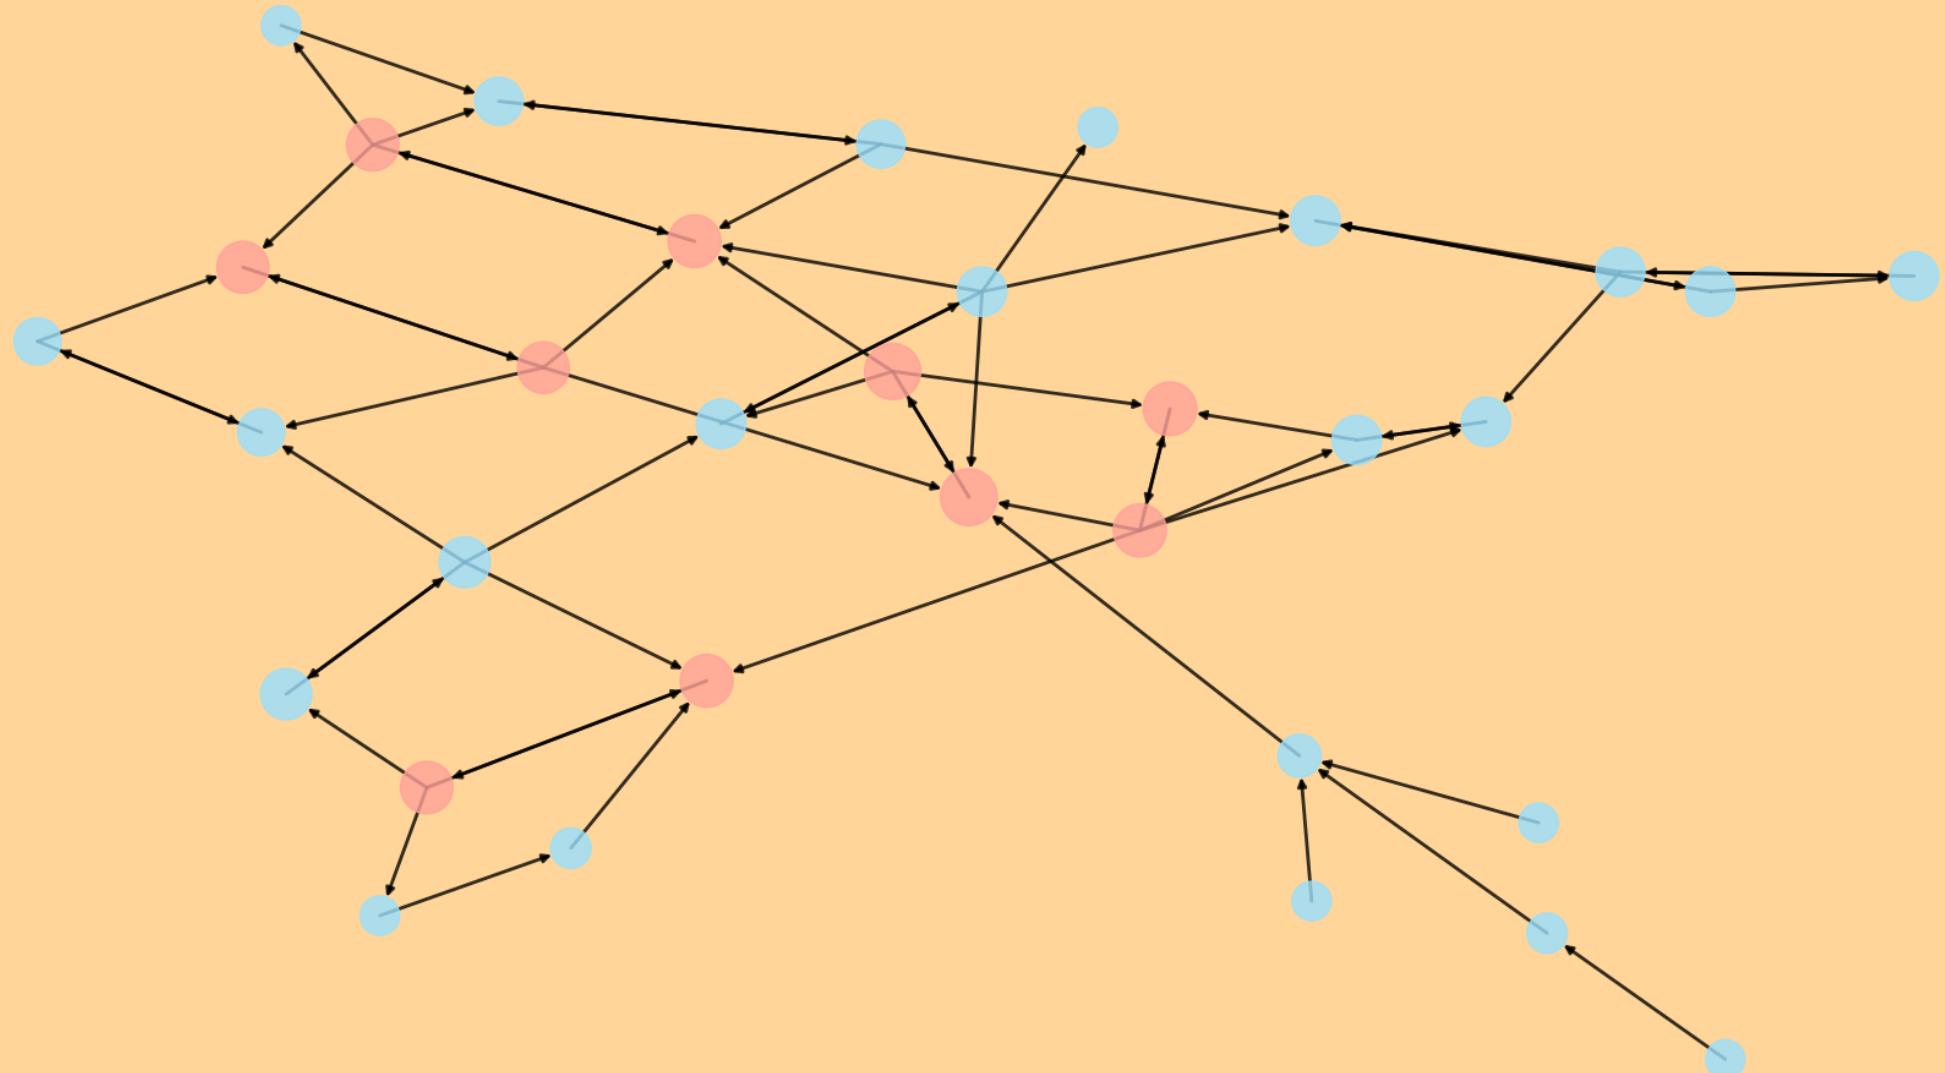


<sup>1</sup>Nemati H., Azhari V., et al., Host-based Virtual Machine Workload Characterization using Hypervisor Trace Mining, ACM Trans 2019

# Use case II

---

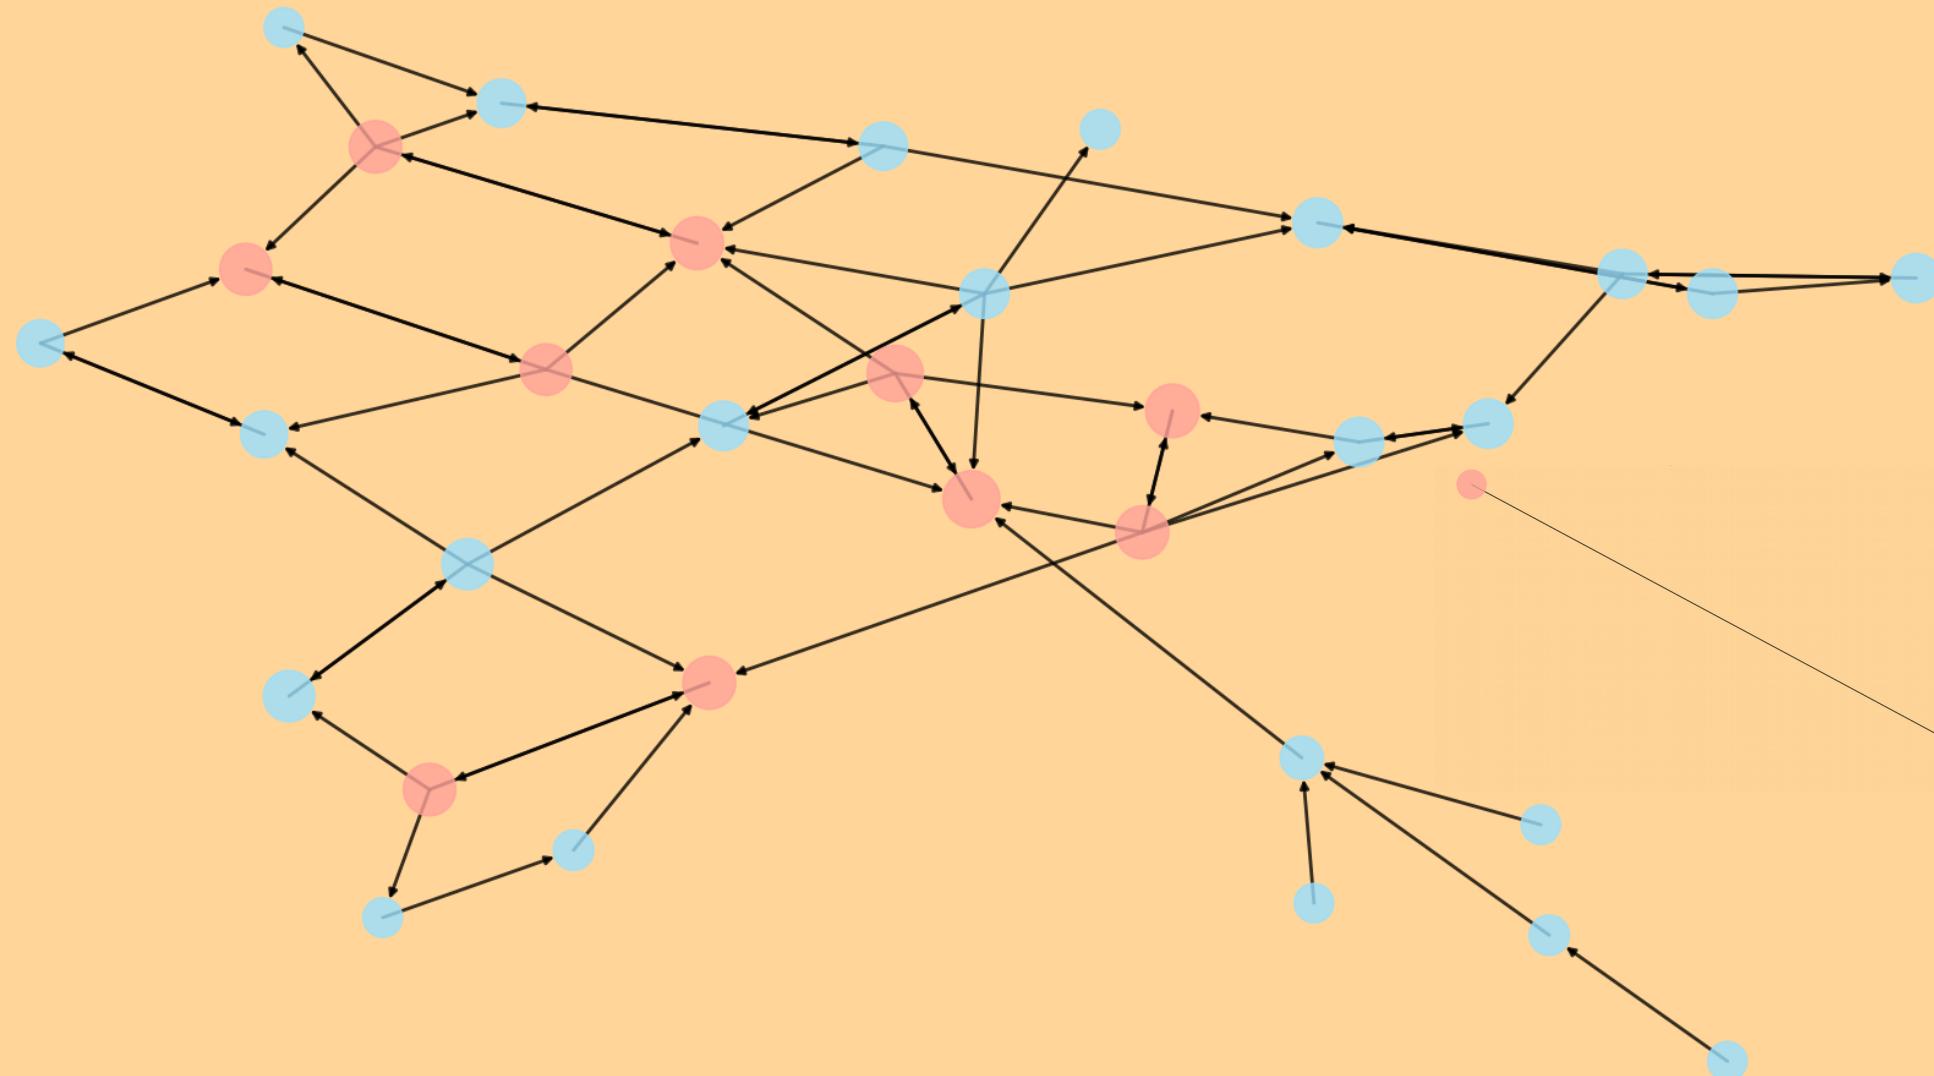
Subgraph I:



# Use case II

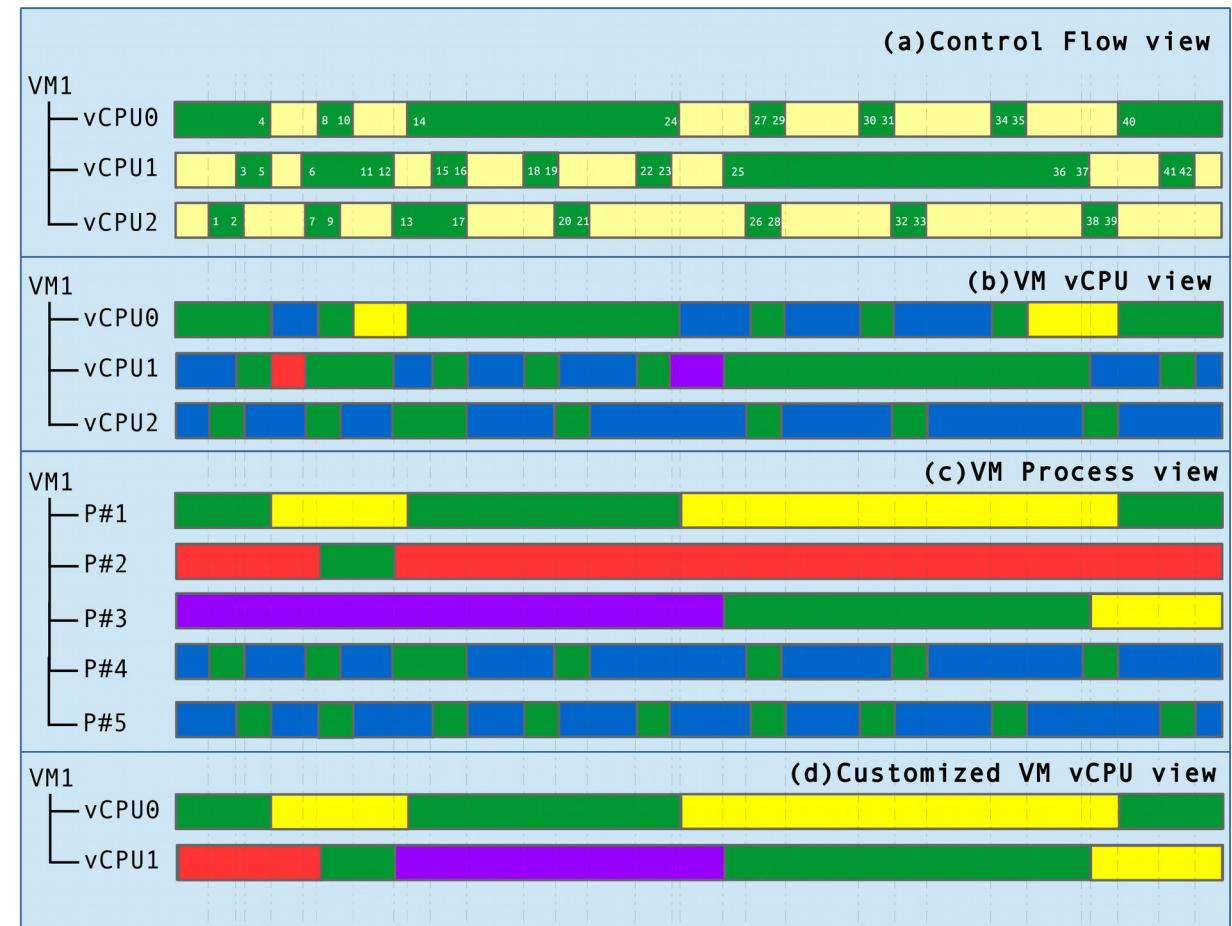
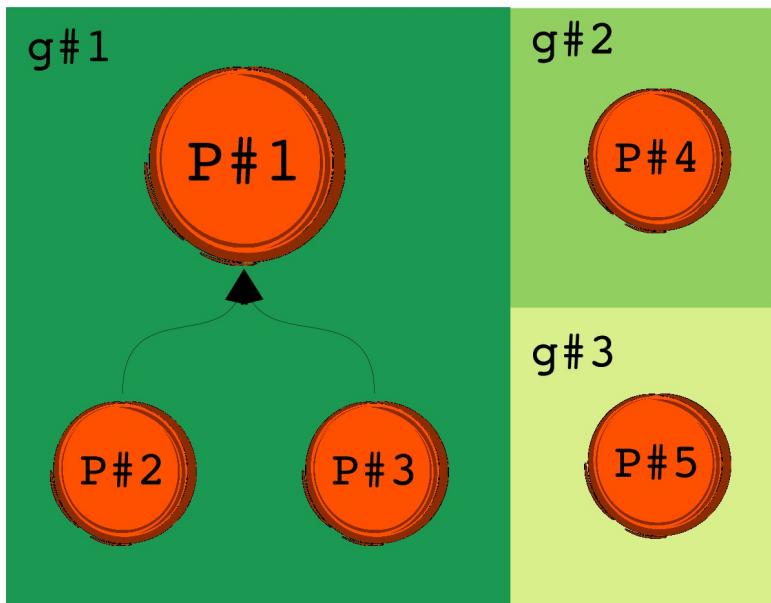
---

Subgraph I and II:



# Use case II

- ✓ A customized vCPU view based on active processes
- ✓ Example:

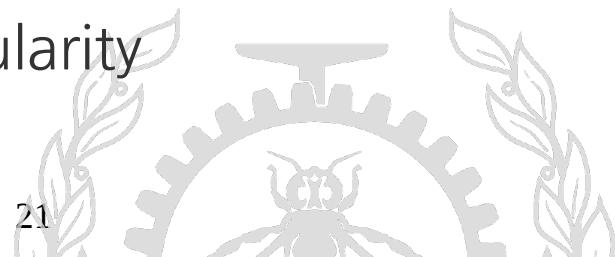


# Conclusions & Future work

---



- Filtering based on process ranking
- Top interacting processes in one system
- Interaction between VM's processes in a cloud environment
- Time-window feature extraction for machine learning applications
- Enhance graph structure analysis, e.g., graph modularity



# Questions

---

*Thanks for your attention!*



# Appendix

# Appendix

---

- VM

---

**Algorithm 2:** Process Ranking Algorithm (PRA)

---

```
1 if event == sched_ttwu then
2     j = getVMvCPU(wakee_tid);
3     k = getVMvCPU(waker_tid);
4     wakerCR3 = getLastCR3(vCPUk);
5     updateWaker(vCPUj, wakerCR3);
6 else if event == vm_inj_virq then
7     j = getVMvCPU(tid);
8     pCR3 = getVMvCPU(vCPUj);
9     if vec == IPI then
10        wakerCR3 = queryWaker(vCPUj);
11        LINK_HORIZONTAL(wakerCR3, pCR3);
12    for all process cr3i ∈ CR3 do
13        R(cr3i);
14    connected_subgraph = breadth-first-search(CR3);
15    customized_graph = denoise(connected_subgraph, R(CR3));
```

---